

Problem 1: Basics of Neural Networks

- **Learning Objective:** In this problem, you are asked to implement a basic multi-layer fully connected neural network from scratch, including forward and backward passes of certain essential layers, to perform an image classification task on the CIFAR100 dataset. You need to implement essential functions in different indicated python files under directory `lib`.
- **Provided Code:** We provide the skeletons of classes you need to complete. Forward checking and gradient checkings are provided for verifying your implementation as well.
- **TODOs:** You are asked to implement the forward passes and backward passes for standard layers and loss functions, various widely-used optimizers, and part of the training procedure. And finally we want you to train a network from scratch on your own. Also, there are inline questions you need to answer. See `README.md` to set up your environment.

```
In [1]: from lib.mlp.fully_conn import *
        from lib.mlp.layer_utils import *
        from lib.datasets import *
        from lib.mlp.train import *
        from lib.grad_check import *
        from lib.optim import *
        import numpy as np
        import matplotlib.pyplot as plt

        %matplotlib inline
        plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
        plt.rcParams['image.interpolation'] = 'nearest'
        plt.rcParams['image.cmap'] = 'gray'

        # for auto-reloading external modules
        # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
        %load_ext autoreload
        %autoreload 2
```

Loading the data (CIFAR-100 with 20 superclasses)

In this homework, we will be classifying images from the CIFAR-100 dataset into the 20 superclasses. More information about the CIFAR-100 dataset and the 20 superclasses can be found [here](#).

Download the CIFAR-100 data files [here](#), and save the `.mat` files to the `data/cifar100` directory.

Load the dataset.

```
In [3]: data = CIFAR100_data('data/cifar100/')
for k, v in data.items():
    if type(v) == np.ndarray:
        print ("Name: {} Shape: {}, {}".format(k, v.shape, type(v)))
    else:
        print("{}: {}".format(k, v))
label_names = data['label_names']
mean_image = data['mean_image'][0]
std_image = data['std_image'][0]
```

```
Name: data_train Shape: (40000, 32, 32, 3), <class 'numpy.ndarray'>
Name: labels_train Shape: (40000,), <class 'numpy.ndarray'>
Name: data_val Shape: (10000, 32, 32, 3), <class 'numpy.ndarray'>
Name: labels_val Shape: (10000,), <class 'numpy.ndarray'>
Name: data_test Shape: (10000, 32, 32, 3), <class 'numpy.ndarray'>
Name: labels_test Shape: (10000,), <class 'numpy.ndarray'>
label_names: ['aquatic_mammals', 'fish', 'flowers', 'food_containers', 'fruit_and_vegetables', 'household_electrical_devices', 'household_furniture', 'insects', 'large_carnivores', 'large_man-made_outdoor_things', 'large_natural_outdoor_scenes', 'large_omnivores_and_herbivores', 'medium_mammals', 'non-insect_invertebrates', 'people', 'reptiles', 'small_mammals', 'trees', 'vehicles_1', 'vehicles_2']
Name: mean_image Shape: (1, 1, 1, 3), <class 'numpy.ndarray'>
Name: std_image Shape: (1, 1, 1, 3), <class 'numpy.ndarray'>
```

Implement Standard Layers

You will now implement all the following standard layers commonly seen in a fully connected neural network (aka multi-layer perceptron, MLP). Please refer to the file `lib/mlp/layer_utils.py`. Take a look at each class skeleton, and we will walk you through the network layer by layer. We provide results of some examples we pre-computed for you for checking the forward pass, and also the gradient checking for the backward pass.

FC Forward [2pt]

In the class skeleton `flatten` and `fc` in `lib/mlp/layer_utils.py`, please complete the forward pass in function `forward`. The input to the `fc` layer may not be of dimension (batch size, features size), it could be an image or any higher dimensional data. We want to convert the input to have a shape of (batch size, features size). Make sure that you handle this dimensionality issue.

```

In [246... %reload_ext autoreload

# Test the fc forward function
input_bz = 3 # batch size
input_dim = (7, 6, 4)
output_dim = 4

input_size = input_bz * np.prod(input_dim)
weight_size = output_dim * np.prod(input_dim)

flatten_layer = flatten(name="flatten_test")
single_fc = fc(np.prod(input_dim), output_dim, init_scale=0.02, name="fc_test")

x = np.linspace(-0.1, 0.4, num=input_size).reshape(input_bz, *input_dim)
w = np.linspace(-0.2, 0.2, num=weight_size).reshape(np.prod(input_dim), output_dim)
b = np.linspace(-0.3, 0.3, num=output_dim)

single_fc.params[single_fc.w_name] = w
single_fc.params[single_fc.b_name] = b

out = single_fc.forward(flatten_layer.forward(x))

correct_out = np.array([[0.63910291, 0.83740057, 1.03569824, 1.23399591],
                        [0.61401587, 0.82903823, 1.04406058, 1.25908294],
                        [0.58892884, 0.82067589, 1.05242293, 1.28416997]])

# Compare your output with the above pre-computed ones.
# The difference should not be larger than 1e-8
print ("Difference: ", rel_error(out, correct_out))

Difference: 4.026016656214849e-09

```

FC Backward [2pt]

Please complete the function `backward` as the backward pass of the `flatten` and `fc` layers. Follow the instructions in the comments to store gradients into the predefined dictionaries in the attributes of the class. Parameters of the layer are also stored in the predefined dictionary.

In [312]...

```
%reload_ext autoreload

# Test the fc backward function
inp = np.random.randn(15, 2, 2, 3)
w = np.random.randn(12, 15)
b = np.random.randn(15)
dout = np.random.randn(15, 15)

flatten_layer = flatten(name="flatten_test")
x = flatten_layer.forward(inp)
single_fc = fc(np.prod(x.shape[1:]), 15, init_scale=5e-2, name="fc_test")
single_fc.params[single_fc.w_name] = w
single_fc.params[single_fc.b_name] = b

dx_num = eval_numerical_gradient_array(lambda x: single_fc.forward(x), x, dout)
dw_num = eval_numerical_gradient_array(lambda w: single_fc.forward(x), w, dout)
db_num = eval_numerical_gradient_array(lambda b: single_fc.forward(x), b, dout)

out = single_fc.forward(x)
dx = single_fc.backward(dout)
dw = single_fc.grads[single_fc.w_name]
db = single_fc.grads[single_fc.b_name]
dinp = flatten_layer.backward(dx)

# The error should be around 1e-9
print("dx Error: ", rel_error(dx_num, dx))
# The errors should be around 1e-10
print("dw Error: ", rel_error(dw_num, dw))
print("db Error: ", rel_error(db_num, db))
# The shapes should be same
print("dinp Shape: ", dinp.shape, inp.shape)

dx Error:  1.9504093719045184e-09
dw Error:  1.305653099063425e-09
db Error:  3.4065475213175096e-10
dinp Shape:  (15, 2, 2, 3) (15, 2, 2, 3)
```

GeLU Forward [2pt]

In the class skeleton `gelu` in `lib/mlp/layer_utils.py`, please complete the forward pass.

GeLU is a smooth version of ReLU and it's used in pre-training LLMs such as GPT-3 and BERT.

$$\text{GeLU}(x) = x\Phi(x) \approx 0.5x(1 + \tanh(\sqrt{2/\pi}(x + 0.044715x^3)))$$

Where $\Phi(x)$ is the CDF for standard Gaussian random variables. You should use the approximate version to compute forward and backward pass.

In [130]...

```
%reload_ext autoreload

# Test the leaky_relu forward function
x = np.linspace(-1.5, 1.5, num=12).reshape(3, 4)
gelu_f = gelu(name="gelu_f")

out = gelu_f.forward(x)
correct_out = np.array([[-0.10042842, -0.13504766, -0.16231757, -0.1689214 ]
                        [-0.13960493, -0.06078651,  0.07557713,  0.26948598]
                        [ 0.51289678,  0.79222788,  1.09222506,  1.39957158]])

# Compare your output with the above pre-computed ones.
# The difference should not be larger than 1e-7
print ("Difference: ", rel_error(out, correct_out))
```

Difference: 1.8037541876132445e-08

GeLU Backward [2pt]

Please complete the `backward` pass of the class `gelu`.

In [131]...

```
%reload_ext autoreload

# Test the relu backward function
x = np.random.randn(15, 15)
dout = np.random.randn(*x.shape)
gelu_b = gelu(name="gelu_b")

dx_num = eval_numerical_gradient_array(lambda x: gelu_b.forward(x), x, dout)

out = gelu_b.forward(x)
dx = gelu_b.backward(dout)

# The error should not be larger than 1e-4, since we are using an approximate
print ("dx Error: ", rel_error(dx_num, dx))
```

dx Error: 9.494869310277931e-06

Dropout Forward [2pt]

In the class `dropout` in `lib/mlp/layer_utils.py`, please complete the `forward` pass.

Remember that the dropout is **only applied during training phase**, you should pay attention to this while implementing the function.

Important Note1: The probability argument input to the function is the "keep probability": probability that each activation is kept.

Important Note2: If the `keep_prob` is set to 1, make it as no dropout.

```
In [174... %reload_ext autoreload

x = np.random.randn(100, 100) + 5.0

print ("-----")
for p in [0, 0.25, 0.50, 0.75, 1]:
    dropout_f = dropout(keep_prob=p)
    out = dropout_f.forward(x, True)
    out_test = dropout_f.forward(x, False)

    # Mean of output should be similar to mean of input
    # Means of output during training time and testing time should be similar
    print ("Dropout Keep Prob = ", p)
    print ("Mean of input: ", x.mean())
    print ("Mean of output during training time: ", out.mean())
    print ("Mean of output during testing time: ", out_test.mean())
    print ("Fraction of output set to zero during training time: ", (out == 0).mean())
    print ("Fraction of output set to zero during testing time: ", (out_test == 0).mean())
    print ("-----")
```

```
-----
Dropout Keep Prob = 0
Mean of input: 4.992070685477767
Mean of output during training time: 4.992070685477767
Mean of output during testing time: 4.992070685477767
Fraction of output set to zero during training time: 0.0
Fraction of output set to zero during testing time: 0.0
-----
```

```
Dropout Keep Prob = 0.25
Mean of input: 4.992070685477767
Mean of output during training time: 4.804921321508981
Mean of output during testing time: 4.992070685477767
Fraction of output set to zero during training time: 0.7594
Fraction of output set to zero during testing time: 0.0
-----
```

```
Dropout Keep Prob = 0.5
Mean of input: 4.992070685477767
Mean of output during training time: 4.940950244843111
Mean of output during testing time: 4.992070685477767
Fraction of output set to zero during training time: 0.5047
Fraction of output set to zero during testing time: 0.0
-----
```

```
Dropout Keep Prob = 0.75
Mean of input: 4.992070685477767
Mean of output during training time: 5.072558378654152
Mean of output during testing time: 4.992070685477767
Fraction of output set to zero during training time: 0.2378
Fraction of output set to zero during testing time: 0.0
-----
```

```
Dropout Keep Prob = 1
Mean of input: 4.992070685477767
Mean of output during training time: 4.992070685477767
Mean of output during testing time: 4.992070685477767
Fraction of output set to zero during training time: 0.0
Fraction of output set to zero during testing time: 0.0
-----
```

Dropout Backward [2pt]

Please complete the backward pass. Again remember that the dropout is only applied during training phase, handle this in the backward pass as well.

```
In [180... %reload_ext autoreload

x = np.random.randn(5, 5) + 5
dout = np.random.randn(*x.shape)

keep_prob = 0.75
dropout_b = dropout(keep_prob, seed=100)
out = dropout_b.forward(x, True, seed=1)
dx = dropout_b.backward(dout)
dx_num = eval_numerical_gradient_array(lambda xx: dropout_b.forward(xx, True

# The error should not be larger than 1e-10
print ('dx relative error: ', rel_error(dx, dx_num))

dx relative error: 3.003117874135711e-11
```

Testing cascaded layers: FC + GeLU [2pt]

Please find the TestFCGeLU function in lib/mlp/fully_conn.py .

You only need to complete a few lines of code in the TODO block.

Please design an Flatten -> FC -> GeLU network where the parameters of them match the given x, w, and b.

Please insert the corresponding names you defined for each layer to param_name_w, and param_name_b respectively. Here you only modify the param_name part, the _w , and _b are automatically assigned during network setup

In [254...

```

%reload_ext autoreload

x = np.random.randn(3, 5, 3) # the input features
w = np.random.randn(15, 5)   # the weight of fc layer
b = np.random.randn(5)       # the bias of fc layer
dout = np.random.randn(3, 5) # the gradients to the output, notice the shape

tiny_net = TestFCGeLU()

#####
# TODO: param_name should be replaced accordingly #
#####
tiny_net.net.assign("fc1_w", w)
tiny_net.net.assign("fc1_b", b)
#####
#                               END OF YOUR CODE                               #
#####

out = tiny_net.forward(x)
dx = tiny_net.backward(dout)

#####
# TODO: param_name should be replaced accordingly #
#####
dw = tiny_net.net.get_grads("fc1_w")
db = tiny_net.net.get_grads("fc1_b")
#####
#                               END OF YOUR CODE                               #
#####

dx_num = eval_numerical_gradient_array(lambda x: tiny_net.forward(x), x, dout)
dw_num = eval_numerical_gradient_array(lambda w: tiny_net.forward(x), w, dout)
db_num = eval_numerical_gradient_array(lambda b: tiny_net.forward(x), b, dout)

# The errors should not be larger than 1e-7
print ("dx error: ", rel_error(dx_num, dx))
print ("dw error: ", rel_error(dw_num, dw))
print ("db error: ", rel_error(db_num, db))

(3, 15)
(3, 5)
(15,)
(5,)
(15, 5)
dx error: 4.8513086385104e-06
dw error: 3.8899350805416096e-05
db error: 3.421977219350853e-05

```


SoftMax Function and Loss Layer [2pt]

In the `lib/mlp/layer_utils.py`, please first complete the function `softmax`, which will be used in the function `cross_entropy`. Then, implement `cross_entropy` using `softmax`. Please refer to the lecture slides of the mathematical expressions of the cross entropy loss function, and complete its forward pass and backward pass. You should also take care of `size_average` on whether or not to divide by the batch size.

```
In [238... %reload_ext autoreload

num_classes, num_inputs = 6, 100
x = 0.001 * np.random.randn(num_inputs, num_classes)
y = np.random.randint(num_classes, size=num_inputs)

test_loss = cross_entropy()

dx_num = eval_numerical_gradient(lambda x: test_loss.forward(x, y), x, verbose=0)

loss = test_loss.forward(x, y)
dx = test_loss.backward()

# Test softmax_loss function. Loss should be around 1.792
# and dx error should be at the scale of 1e-8 (or smaller)
print ("Cross Entropy Loss: ", loss)
print ("dx error: ", rel_error(dx_num, dx))

Cross Entropy Loss:  1.7918038711988697
dx error:  8.063889982552038e-09
```

Test a Small Fully Connected Network [2pt]

Please find the `SmallFullyConnectedNetwork` function in `lib/mlp/fully_conn.py`.

Again you only need to complete few lines of code in the TODO block.

Please design an FC → GeLU → FC network where the shapes of parameters match the given shapes.

Please insert the corresponding names you defined for each layer to `param_name_w`, and `param_name_b` respectively.

Here you only modify the `param_name` part, the `_w`, and `_b` are automatically assigned during network setup.

In [315...

```

%reload_ext autoreload

seed = 1234
np.random.seed(seed=seed)

model = SmallFullyConnectedNetwork()
loss_func = cross_entropy()

N, D, = 4, 4 # N: batch size, D: input dimension
H, C = 30, 7 # H: hidden dimension, C: output dimension
std = 0.02
x = np.random.randn(N, D)
y = np.random.randint(C, size=N)

print ("Testing initialization ... ")

#####
# TODO: param_name should be replaced accordingly #
#####
w1_std = abs(model.net.get_params("fc1_w").std() - std)
b1 = model.net.get_params("fc1_b").std()
w2_std = abs(model.net.get_params("fc2_w").std() - std)
b2 = model.net.get_params("fc2_b").std()
#####
#                               END OF YOUR CODE                               #
#####

assert w1_std < std / 10, "First layer weights do not seem right"
assert np.all(b1 == 0), "First layer biases do not seem right"
assert w2_std < std / 10, "Second layer weights do not seem right"
assert np.all(b2 == 0), "Second layer biases do not seem right"
print ("Passed!")

print ("Testing test-time forward pass ... ")
w1 = np.linspace(-0.7, 0.3, num=D*H).reshape(D, H)
w2 = np.linspace(-0.2, 0.2, num=H*C).reshape(H, C)
b1 = np.linspace(-0.6, 0.2, num=H)
b2 = np.linspace(-0.9, 0.1, num=C)

#####
# TODO: param_name should be replaced accordingly #
#####
model.net.assign("fc1_w", w1)
model.net.assign("fc1_b", b1)
model.net.assign("fc2_w", w2)
model.net.assign("fc2_b", b2)
#####
#                               END OF YOUR CODE                               #
#####

feats = np.linspace(-5.5, 4.5, num=N*D).reshape(D, N).T
scores = model.forward(feats)
correct_scores = np.asarray([[-2.33881897, -1.92174121, -1.50466344, -1.0875
                             [-1.57214916, -1.1857013 , -0.79925345, -0.4128
                             [-0.80178618, -0.44604469, -0.0903032 ,  0.2654
                             [-0.00331319,  0.32124836,  0.64580991,  0.9703

```

```

scores_diff = np.sum(np.abs(scores - correct_scores))
assert scores_diff < 1e-6, "Your implementation might be wrong!"
print ("Passed!")

print ("Testing the loss ...",)
y = np.asarray([0, 5, 1, 4])
loss = loss_func.forward(scores, y)
dLoss = loss_func.backward()
correct_loss = 2.4248995879903195
assert abs(loss - correct_loss) < 1e-10, "Your implementation might be wrong"
print ("Passed!")

print ("Testing the gradients (error should be no larger than 1e-6) ...")
din = model.backward(dLoss)
for layer in model.net.layers:
    if not layer.params:
        continue
    for name in sorted(layer.grads):
        f = lambda _: loss_func.forward(model.forward(feats), y)
        grad_num = eval_numerical_gradient(f, layer.params[name], verbose=False)
        print ('%s relative error: %.2e' % (name, rel_error(grad_num, layer.

```

Testing initialization ...

Passed!

Testing test-time forward pass ...

Passed!

Testing the loss ...

Passed!

Testing the gradients (error should be no larger than 1e-6) ...

fc1_b relative error: 1.31e-08

fc1_w relative error: 3.53e-08

fc2_b relative error: 4.01e-10

fc2_w relative error: 2.50e-08

Test a Fully Connected Network regularized with Dropout [2pt]

Please find the `DropoutNet` function in `fully_conn.py` under `lib/mlp` directory.

For this part you don't need to design a new network, just simply run the following test code.

If something goes wrong, you might want to double check your dropout implementation.

In [321]...

```
%reload_ext autoreload

seed = 1234
np.random.seed(seed=seed)

N, D, C = 3, 15, 10
X = np.random.randn(N, D)
y = np.random.randint(C, size=(N,))

for keep_prob in [0, 0.25, 0.5]:
    np.random.seed(seed=seed)
    print ("Dropout p =", keep_prob)
    model = DropoutNet(keep_prob=keep_prob, seed=seed)
    loss_func = cross_entropy()
    output = model.forward(X, True, seed=seed)
    loss = loss_func.forward(output, y)
    dLoss = loss_func.backward()
    dX = model.backward(dLoss)
    grads = model.net.grads

    print ("Error of gradients should be around or less than 1e-3")
    for name in sorted(grads):
        if name not in model.net.params.keys():
            continue
        f = lambda _: loss_func.forward(model.forward(X, True, seed=seed), y)
        grad_num = eval_numerical_gradient(f, model.net.params[name], verbose=0)
        print ("{} relative error: {}".format(name, rel_error(grad_num, grads[name])))
    print ()
```

Dropout p = 0

Error of gradients should be around or less than 1e-3

fc1_b relative error: 2.8516549520063804e-07

fc1_w relative error: 3.7626907492775348e-06

fc2_b relative error: 1.3390330536574157e-08

fc2_w relative error: 3.0874875391929026e-05

fc3_b relative error: 3.171994103778982e-10

fc3_w relative error: 2.0488862038376876e-06

Dropout p = 0.25

Error of gradients should be around or less than 1e-3

fc1_b relative error: 3.2230322968934954e-07

fc1_w relative error: 2.784401997035636e-06

fc2_b relative error: 1.4909849738002271e-07

fc2_w relative error: 3.137678111642197e-05

fc3_b relative error: 6.679255248099083e-11

fc3_w relative error: 4.7620800542288264e-07

Dropout p = 0.5

Error of gradients should be around or less than 1e-3

fc1_b relative error: 9.415776886825045e-07

fc1_w relative error: 1.0482378081948802e-06

fc2_b relative error: 1.3366658240212987e-08

fc2_w relative error: 8.895735135607248e-06

fc3_b relative error: 4.395930640345582e-10

fc3_w relative error: 4.179186917341344e-06

Training a Network

In this section, we defined a `TinyNet` class for you to fill in the TODO block in `lib/mlp/fully_conn.py`.

- Here please design a two layer fully connected network with Leaky ReLU activation (Flatten --> FC --> GeLU --> FC).
- You can adjust the number of hidden neurons, `batch_size`, `epochs`, and learning rate decay parameters.
- Please read the `lib/train.py` carefully and complete the TODO blocks in the `train_net` function first. Codes in "Test a Small Fully Connected Network" can be helpful.
- Implement SGD in `lib/optim.py`, you will be asked to complete weight decay and Adam in the later sections.

```
In [317... # Arrange the data
data_dict = {
    "data_train": (data["data_train"], data["labels_train"]),
    "data_val": (data["data_val"], data["labels_val"]),
    "data_test": (data["data_test"], data["labels_test"])
}
```

```
In [318... print("Data shape:", data["data_train"].shape)
print("Flattened data input size:", np.prod(data["data_train"].shape[1:]))
print("Number of data classes:", max(data['labels_train']) + 1)
```

```
Data shape: (40000, 32, 32, 3)
Flattened data input size: 3072
Number of data classes: 20
```

Now train the network to achieve at least 30% validation accuracy [5pt]

You may only adjust the hyperparameters inside the TODO block

```
In [319... %autoreload
```

In [333]...

```
%reload_ext autoreload

seed = 123
np.random.seed(seed=seed)

model = TinyNet()
loss_f = cross_entropy()
optimizer = SGD(model.net, 0.1)

results = None
#####
# TODO: Use the train_net function you completed to train a network
#####

batch_size = 100
epochs = 5
lr_decay = 0.99
lr_decay_every = 100

#####
#                                     END OF YOUR CODE
#####
results = train_net(data_dict, model, loss_f, optimizer, batch_size, epochs,
                    lr_decay, lr_decay_every, show_every=10000, verbose=True)
opt_params, loss_hist, train_acc_hist, val_acc_hist = results
```

```
2%|██████████|
| 7/400 [00:00<00:11, 34.32it/s]
(Iteration 1 / 2000) Average loss: 2.995905928014045

100%|██████████| 400/400 [00:11
<00:00, 35.98it/s]
(Epoch 1 / 5) Training Accuracy: 0.2856, Validation Accuracy: 0.2762

100%|██████████| 400/400 [00:10
<00:00, 37.39it/s]
(Epoch 2 / 5) Training Accuracy: 0.31345, Validation Accuracy: 0.2843

100%|██████████| 400/400 [00:11
<00:00, 34.19it/s]
(Epoch 3 / 5) Training Accuracy: 0.352625, Validation Accuracy: 0.3091

100%|██████████| 400/400 [00:10
<00:00, 36.61it/s]
(Epoch 4 / 5) Training Accuracy: 0.370725, Validation Accuracy: 0.3078

100%|██████████| 400/400 [00:10
<00:00, 39.43it/s]
(Epoch 5 / 5) Training Accuracy: 0.403625, Validation Accuracy: 0.3243
```

In [334]...

```
# Take a look at what names of params were stored
print (opt_params.keys())


dict_keys(['fc1_w', 'fc1_b', 'fc2_w', 'fc2_b'])
```

```
In [335... # Demo: How to load the parameters to a newly defined network
model = TinyNet()
model.net.load(opt_params)
val_acc = compute_acc(model, data["data_val"], data["labels_val"])
print ("Validation Accuracy: {}".format(val_acc*100))
test_acc = compute_acc(model, data["data_test"], data["labels_test"])
print ("Testing Accuracy: {}".format(test_acc*100))
```

```
Loading Params: fc1_w Shape: (3072, 300)
Loading Params: fc1_b Shape: (300,)
Loading Params: fc2_w Shape: (300, 20)
Loading Params: fc2_b Shape: (20,)
Validation Accuracy: 32.43%
Testing Accuracy: 32.09%
```

```
In [336... # Plot the learning curves
plt.subplot(2, 1, 1)
plt.title('Training loss')
loss_hist_ = loss_hist[1::100] # sparse the curve a bit
plt.plot(loss_hist_, '-o')
plt.xlabel('Iteration')

plt.subplot(2, 1, 2)
plt.title('Accuracy')
plt.plot(train_acc_hist, '-o', label='Training')
plt.plot(val_acc_hist, '-o', label='Validation')
plt.xlabel('Epoch')
plt.legend(loc='lower right')
plt.gcf().set_size_inches(15, 12)
plt.show()
```



Different Optimizers and Regularization Techniques

There are several more advanced optimizers than vanilla SGD, and there are many regularization tricks. You'll implement them in this section. Please complete the TODOs in the `lib/optim.py`.

SGD + Weight Decay [2pt]

The update rule of SGD plus weight decay is as shown below:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} J(\theta_t) - \lambda \theta_t$$

Update the `SGD()` function in `lib/optim.py`, and also incorporate weight decay options.


```

In [387... %reload_ext autoreload

# Test the implementation of SGD with Momentum
seed = 1234
np.random.seed(seed=seed)

N, D = 4, 5
test_sgd = sequential(fc(N, D, name="sgd_fc"))

w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)

test_sgd.layers[0].params = {"sgd_fc_w": w}
test_sgd.layers[0].grads = {"sgd_fc_w": dw}

test_sgd_wd = SGD(test_sgd, 1e-3, 1e-4)
test_sgd_wd.step()

updated_w = test_sgd.layers[0].params["sgd_fc_w"]

expected_updated_w = np.asarray([
    [-0.39936, -0.34678632, -0.29421263, -0.24163895, -0.18906526],
    [-0.13649158, -0.08391789, -0.03134421, 0.02122947, 0.07380316],
    [ 0.12637684, 0.17895053, 0.23152421, 0.28409789, 0.33667158],
    [ 0.38924526, 0.44181895, 0.49439263, 0.54696632, 0.59954 ]])

print ('The following errors should be around or less than 1e-6')
print ('updated_w error: ', rel_error(updated_w, expected_updated_w))

The following errors should be around or less than 1e-6
updated_w error: 8.677112905190533e-08

```

Comparing SGD and SGD with Weight Decay [2pt]

Run the following code block to train a multi-layer fully connected network with both SGD and SGD plus Weight Decay. You are expected to see Weight Decay have better validation accuracy than vanilla SGD.

```

In [388... seed = 1234

# Arrange a small data
num_train = 20000
small_data_dict = {
    "data_train": (data["data_train"][:num_train], data["labels_train"][:num_train]),
    "data_val": (data["data_val"], data["labels_val"]),
    "data_test": (data["data_test"], data["labels_test"])
}

reset_seed(seed=seed)
model_sgd = FullyConnectedNetwork()
loss_f_sgd = cross_entropy()
optimizer_sgd = SGD(model_sgd.net, 0.01)
print("Training with Vanilla SGD...")
results_sgd = train_net(small_data_dict, model_sgd, loss_f_sgd, optimizer_sgd,
                        max_epochs=50, show_every=10000, verbose=True)

reset_seed(seed=seed)
model_sgdw = FullyConnectedNetwork()
loss_f_sgdw = cross_entropy()
optimizer_sgdw = SGD(model_sgdw.net, 0.01, 1e-4)
print("\nTraining with SGD plus Weight Decay...")
results_sgdw = train_net(small_data_dict, model_sgdw, loss_f_sgdw, optimizer_sgdw,
                        max_epochs=50, show_every=10000, verbose=True)

opt_params_sgd, loss_hist_sgd, train_acc_hist_sgd, val_acc_hist_sgd = results_sgd
opt_params_sgdw, loss_hist_sgdw, train_acc_hist_sgdw, val_acc_hist_sgdw = results_sgdw

plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd, 'o', label="Vanilla SGD")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd, '-o', label="Vanilla SGD")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd, '-o', label="Vanilla SGD")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgdw, 'o', label="SGD with Weight Decay")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgdw, '-o', label="SGD with Weight Decay")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgdw, '-o', label="SGD with Weight Decay")

for i in [1, 2, 3]:

```

```
plt.subplot(3, 1, i)
plt.legend(loc='upper center', ncol=4)
plt.gcf().set_size_inches(15, 15)
plt.show()
```

Training with Vanilla SGD...

```
2%|██████████| 4/200 [00:00<00:05, 36.33it/s]
(Iteration 1 / 10000) Average loss: 3.3332154539088985

100%|██████████| 200/200 [00:05<00:00, 38.18it/s]
(Epoch 1 / 50) Training Accuracy: 0.15095, Validation Accuracy: 0.1474

100%|██████████| 200/200 [00:05<00:00, 38.79it/s]
(Epoch 2 / 50) Training Accuracy: 0.18815, Validation Accuracy: 0.1805

100%|██████████| 200/200 [00:04<00:00, 43.05it/s]
(Epoch 3 / 50) Training Accuracy: 0.2107, Validation Accuracy: 0.2029

100%|██████████| 200/200 [00:04<00:00, 43.72it/s]
(Epoch 4 / 50) Training Accuracy: 0.2314, Validation Accuracy: 0.212

100%|██████████| 200/200 [00:04<00:00, 40.66it/s]
(Epoch 5 / 50) Training Accuracy: 0.23915, Validation Accuracy: 0.2197

100%|██████████| 200/200 [00:04<00:00, 40.94it/s]
(Epoch 6 / 50) Training Accuracy: 0.2552, Validation Accuracy: 0.2298

100%|██████████| 200/200 [00:05<00:00, 39.87it/s]
(Epoch 7 / 50) Training Accuracy: 0.26645, Validation Accuracy: 0.2403

100%|██████████| 200/200 [00:05<00:00, 38.36it/s]
(Epoch 8 / 50) Training Accuracy: 0.27555, Validation Accuracy: 0.2414

100%|██████████| 200/200 [00:04<00:00, 41.67it/s]
(Epoch 9 / 50) Training Accuracy: 0.28185, Validation Accuracy: 0.2413

100%|██████████| 200/200 [00:05<00:00, 38.99it/s]
(Epoch 10 / 50) Training Accuracy: 0.2944, Validation Accuracy: 0.252

100%|██████████| 200/200 [00:05<00:00, 38.98it/s]
(Epoch 11 / 50) Training Accuracy: 0.29735, Validation Accuracy: 0.2543
```

```

100%|██████████| 200/200 [00:05<00:00, 37.02it/s]
(Epoch 12 / 50) Training Accuracy: 0.3021, Validation Accuracy: 0.2587
100%|██████████| 200/200 [00:05<00:00, 39.30it/s]
(Epoch 13 / 50) Training Accuracy: 0.31105, Validation Accuracy: 0.2641
100%|██████████| 200/200 [00:05<00:00, 39.42it/s]
(Epoch 14 / 50) Training Accuracy: 0.3168, Validation Accuracy: 0.2653
100%|██████████| 200/200 [00:04<00:00, 40.24it/s]
(Epoch 15 / 50) Training Accuracy: 0.3217, Validation Accuracy: 0.2681
100%|██████████| 200/200 [00:05<00:00, 36.73it/s]
(Epoch 16 / 50) Training Accuracy: 0.3307, Validation Accuracy: 0.2699
100%|██████████| 200/200 [00:04<00:00, 42.35it/s]
(Epoch 17 / 50) Training Accuracy: 0.33835, Validation Accuracy: 0.2696
100%|██████████| 200/200 [00:05<00:00, 37.66it/s]
(Epoch 18 / 50) Training Accuracy: 0.34565, Validation Accuracy: 0.2737
100%|██████████| 200/200 [00:05<00:00, 39.06it/s]
(Epoch 19 / 50) Training Accuracy: 0.3495, Validation Accuracy: 0.2729
100%|██████████| 200/200 [00:04<00:00, 43.14it/s]
(Epoch 20 / 50) Training Accuracy: 0.35565, Validation Accuracy: 0.2758
100%|██████████| 200/200 [00:04<00:00, 40.17it/s]
(Epoch 21 / 50) Training Accuracy: 0.35825, Validation Accuracy: 0.2729
100%|██████████| 200/200 [00:04<00:00, 43.70it/s]
(Epoch 22 / 50) Training Accuracy: 0.36895, Validation Accuracy: 0.278
100%|██████████| 200/200 [00:04<00:00, 44.87it/s]
(Epoch 23 / 50) Training Accuracy: 0.3734, Validation Accuracy: 0.2783
100%|██████████| 200/200 [00:04<00:00, 44.39it/s]
(Epoch 24 / 50) Training Accuracy: 0.3756, Validation Accuracy: 0.2768

```

```

100%|██████████| 200/200 [00:04<00:00, 44.21it/s]
(Epoch 25 / 50) Training Accuracy: 0.38495, Validation Accuracy: 0.278
100%|██████████| 200/200 [00:04<00:00, 43.90it/s]
(Epoch 26 / 50) Training Accuracy: 0.38415, Validation Accuracy: 0.2757
100%|██████████| 200/200 [00:04<00:00, 45.72it/s]
(Epoch 27 / 50) Training Accuracy: 0.40365, Validation Accuracy: 0.2804
100%|██████████| 200/200 [00:04<00:00, 44.90it/s]
(Epoch 28 / 50) Training Accuracy: 0.40105, Validation Accuracy: 0.2812
100%|██████████| 200/200 [00:04<00:00, 42.99it/s]
(Epoch 29 / 50) Training Accuracy: 0.40885, Validation Accuracy: 0.2773
100%|██████████| 200/200 [00:04<00:00, 44.80it/s]
(Epoch 30 / 50) Training Accuracy: 0.4163, Validation Accuracy: 0.2803
100%|██████████| 200/200 [00:04<00:00, 43.43it/s]
(Epoch 31 / 50) Training Accuracy: 0.41745, Validation Accuracy: 0.2838
100%|██████████| 200/200 [00:04<00:00, 43.89it/s]
(Epoch 32 / 50) Training Accuracy: 0.42125, Validation Accuracy: 0.2758
100%|██████████| 200/200 [00:04<00:00, 45.21it/s]
(Epoch 33 / 50) Training Accuracy: 0.433, Validation Accuracy: 0.2777
100%|██████████| 200/200 [00:04<00:00, 44.35it/s]
(Epoch 34 / 50) Training Accuracy: 0.4322, Validation Accuracy: 0.2782
100%|██████████| 200/200 [00:04<00:00, 45.04it/s]
(Epoch 35 / 50) Training Accuracy: 0.44095, Validation Accuracy: 0.2753
100%|██████████| 200/200 [00:05<00:00, 36.53it/s]
(Epoch 36 / 50) Training Accuracy: 0.4517, Validation Accuracy: 0.2783
100%|██████████| 200/200 [00:05<00:00, 38.36it/s]
(Epoch 37 / 50) Training Accuracy: 0.4583, Validation Accuracy: 0.2759

```

[illegible]

(Epoch 11 / 50) Training Accuracy: 0.2792, Validation Accuracy: 0.2583

```


100%|███████████████████████████████| 200/200 [00:04<00:00, 44.89it/s]  
(Epoch 12 / 50) Training Accuracy: 0.28575, Validation Accuracy: 0.2646



100%|███████████████████████████████| 200/200 [00:04<00:00, 42.03it/s]  
(Epoch 13 / 50) Training Accuracy: 0.2879, Validation Accuracy: 0.2657



100%|███████████████████████████████| 200/200 [00:04<00:00, 44.53it/s]  
(Epoch 14 / 50) Training Accuracy: 0.28865, Validation Accuracy: 0.2664



100%|███████████████████████████████| 200/200 [00:04<00:00, 43.64it/s]  
(Epoch 15 / 50) Training Accuracy: 0.29545, Validation Accuracy: 0.2705



100%|███████████████████████████████| 200/200 [00:04<00:00, 45.96it/s]  
(Epoch 16 / 50) Training Accuracy: 0.2964, Validation Accuracy: 0.2737



100%|███████████████████████████████| 200/200 [00:04<00:00, 43.23it/s]  
(Epoch 17 / 50) Training Accuracy: 0.30345, Validation Accuracy: 0.2752



100%|███████████████████████████████| 200/200 [00:04<00:00, 44.38it/s]  
(Epoch 18 / 50) Training Accuracy: 0.30555, Validation Accuracy: 0.276



100%|███████████████████████████████| 200/200 [00:05<00:00, 39.58it/s]  
(Epoch 19 / 50) Training Accuracy: 0.30715, Validation Accuracy: 0.2821



100%|███████████████████████████████| 200/200 [00:04<00:00, 40.10it/s]  
(Epoch 20 / 50) Training Accuracy: 0.31265, Validation Accuracy: 0.2799



100%|███████████████████████████████| 200/200 [00:05<00:00, 37.24it/s]  
(Epoch 21 / 50) Training Accuracy: 0.31315, Validation Accuracy: 0.2787



100%|███████████████████████████████| 200/200 [00:04<00:00, 45.16it/s]  
(Epoch 22 / 50) Training Accuracy: 0.31755, Validation Accuracy: 0.2836



100%|███████████████████████████████| 200/200 [00:04<00:00, 43.65it/s]  
(Epoch 23 / 50) Training Accuracy: 0.3192, Validation Accuracy: 0.2833



100%|███████████████████████████████| 200/200 [00:04<00:00, 44.69it/s]  
(Epoch 24 / 50) Training Accuracy: 0.31905, Validation Accuracy: 0.2837


```



```
100%|██████████| 200/200 [00:04<00:00, 45.24it/s]
(Epoch 25 / 50) Training Accuracy: 0.32525, Validation Accuracy: 0.2894
100%|██████████| 200/200 [00:04<00:00, 43.91it/s]
(Epoch 26 / 50) Training Accuracy: 0.3238, Validation Accuracy: 0.2895
100%|██████████| 200/200 [00:04<00:00, 43.99it/s]
(Epoch 27 / 50) Training Accuracy: 0.33645, Validation Accuracy: 0.2944
100%|██████████| 200/200 [00:04<00:00, 44.88it/s]
(Epoch 28 / 50) Training Accuracy: 0.33645, Validation Accuracy: 0.2941
100%|██████████| 200/200 [00:05<00:00, 38.44it/s]
(Epoch 29 / 50) Training Accuracy: 0.33695, Validation Accuracy: 0.2953
100%|██████████| 200/200 [00:04<00:00, 40.46it/s]
(Epoch 30 / 50) Training Accuracy: 0.3425, Validation Accuracy: 0.3
100%|██████████| 200/200 [00:04<00:00, 43.56it/s]
(Epoch 31 / 50) Training Accuracy: 0.3406, Validation Accuracy: 0.2982
100%|██████████| 200/200 [00:04<00:00, 44.46it/s]
(Epoch 32 / 50) Training Accuracy: 0.34505, Validation Accuracy: 0.2949
100%|██████████| 200/200 [00:04<00:00, 44.13it/s]
(Epoch 33 / 50) Training Accuracy: 0.34595, Validation Accuracy: 0.3011
100%|██████████| 200/200 [00:04<00:00, 44.72it/s]
(Epoch 34 / 50) Training Accuracy: 0.34755, Validation Accuracy: 0.301
100%|██████████| 200/200 [00:04<00:00, 43.76it/s]
(Epoch 35 / 50) Training Accuracy: 0.3548, Validation Accuracy: 0.3012
100%|██████████| 200/200 [00:04<00:00, 48.02it/s]
(Epoch 36 / 50) Training Accuracy: 0.3552, Validation Accuracy: 0.2995
100%|██████████| 200/200 [00:05<00:00, 39.97it/s]
(Epoch 37 / 50) Training Accuracy: 0.35525, Validation Accuracy: 0.3034
```

```
100%|██████████| 200/200 [00:04<00:00, 44.58it/s]
(EPOCH 38 / 50) Training Accuracy: 0.3593, Validation Accuracy: 0.3017

100%|██████████| 200/200 [00:04<00:00, 43.02it/s]
(EPOCH 39 / 50) Training Accuracy: 0.3648, Validation Accuracy: 0.3048

100%|██████████| 200/200 [00:04<00:00, 44.06it/s]
(EPOCH 40 / 50) Training Accuracy: 0.36665, Validation Accuracy: 0.311

100%|██████████| 200/200 [00:04<00:00, 45.64it/s]
(EPOCH 41 / 50) Training Accuracy: 0.35765, Validation Accuracy: 0.3068

100%|██████████| 200/200 [00:04<00:00, 44.82it/s]
(EPOCH 42 / 50) Training Accuracy: 0.36375, Validation Accuracy: 0.302

100%|██████████| 200/200 [00:04<00:00, 45.13it/s]
(EPOCH 43 / 50) Training Accuracy: 0.3702, Validation Accuracy: 0.3062

100%|██████████| 200/200 [00:04<00:00, 44.98it/s]
(EPOCH 44 / 50) Training Accuracy: 0.37215, Validation Accuracy: 0.306

100%|██████████| 200/200 [00:04<00:00, 45.08it/s]
(EPOCH 45 / 50) Training Accuracy: 0.37475, Validation Accuracy: 0.3037


100%|██████████| 200/200 [00:04<00:00, 44.59it/s]
(EPOCH 46 / 50) Training Accuracy: 0.37205, Validation Accuracy: 0.3089

100%|██████████| 200/200 [00:04<00:00, 42.72it/s]
(EPOCH 47 / 50) Training Accuracy: 0.3827, Validation Accuracy: 0.3097

100%|██████████| 200/200 [00:04<00:00, 43.65it/s]
(EPOCH 48 / 50) Training Accuracy: 0.38395, Validation Accuracy: 0.313

100%|██████████| 200/200 [00:04<00:00, 43.06it/s]
(EPOCH 49 / 50) Training Accuracy: 0.38155, Validation Accuracy: 0.3131

100%|██████████| 200/200 [00:04<00:00, 42.82it/s]
(EPOCH 50 / 50) Training Accuracy: 0.38415, Validation Accuracy: 0.3121
```



SGD with L1 Regularization [2pts]

With L1 Regularization, your regularized loss becomes $\tilde{J}_{\ell_1}(\theta)$ and it's defined as

$$\tilde{J}_{\ell_1}(\theta) = J(\theta) + \lambda \|\theta\|_{\ell_1}$$

where

$$\|\theta\|_{\ell_1} = \sum_{l=1}^n \sum_{k=1}^{n_l} |\theta_{l,k}|$$

Please implment TODO block of `apply_l1_regularization` in `lib/layer_utils`. Such regularization functionality is called after gradient gathering in the `backward` process.

```

In [376... reset_seed(seed=seed)
model_sgd_l1 = FullyConnectedNetwork()
loss_f_sgd_l1 = cross_entropy()
optimizer_sgd_l1 = SGD(model_sgd_l1.net, 0.01)

print ("\nTraining with SGD plus L1 Regularization...")
results_sgd_l1 = train_net(small_data_dict, model_sgd_l1, loss_f_sgd_l1, opt
                           max_epochs=50, show_every=10000, verbose=True, regu

opt_params_sgd_l1, loss_hist_sgd_l1, train_acc_hist_sgd_l1, val_acc_hist_sgd

plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd, 'o', label="Vanilla SGD")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd, '-o', label="Vanilla SGD")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd, '-o', label="Vanilla SGD")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd_l1, 'o', label="SGD with L1 Regularization")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd_l1, '-o', label="SGD with L1 Regularization")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd_l1, '-o', label="SGD with L1 Regularization")

for i in [1, 2, 3]:
    plt.subplot(3, 1, i)
    plt.legend(loc='upper center', ncol=4)
plt.gcf().set_size_inches(15, 15)
plt.show()

```

Training with SGD plus L1 Regularization...

```

0%|
| 0/200 [00:00<?, ?it/s]/Users/chung/Documents/spring_2023/DeepLearningCS56
6/csci566-assignment1/csci566-assignment1/lib/mlp/layer_utils.py:304: Runti
meWarning: overflow encountered in cosh
    term3 = np.cosh(0.0356774*feat_3 + 0.797885*feat)**-2
2%|
| 4/200 [00:00<00:05, 33.19it/s]
(Iteration 1 / 10000) Average loss: 3.3332154539088985
100%|
| 200/200 [00:04
<00:00, 40.77it/s]


```

[illegible]

[illegible]

[illegible]

[illegible]



SGD with L2 Regularization [2pts]

With L2 Regularization, your regularized loss becomes $\tilde{J}_{\ell_2}(\theta)$ and it's defined as

$$\tilde{J}_{\ell_2}(\theta) = J(\theta) + \lambda \|\theta\|_{\ell_2}^2$$

where

$$\|\theta\|_{\ell_2}^2 = \sum_{l=1}^n \sum_{k=1}^{n_l} \theta_{l,k}^2$$

Similarly, implment TODO block of `apply_l2_regularization` in

`lib/layer_utils` . For SGD, you're also asked to find the λ for L2 Regularization such that it achives the EXACTLY SAME effect as weight decay in the previous cells. As a reminder, learning rate is the same as previously, and the weight decay paramter was 1e-4.

```

In [389... reset_seed(seed=seed)
model_sgd_l2 = FullyConnectedNetwork()
loss_f_sgd_l2 = cross_entropy()
optimizer_sgd_l2 = SGD(model_sgd_l2.net, 0.01)
#####
#### Find lambda for L2 regularization so that #####
#### it achieves EXACTLY THE SAME learning curve as weight decay ####
'''
Since the gradient of updates for L2 is  $lr \cdot (2 \cdot \lambda \cdot w)$  and
weight decay updates gradients by  $decay \cdot w$ , we just need to find a lambda whe
learning rate * 2 * lambda = decay
 $0.01 \cdot 2 \cdot \lambda = 1e-4$ 
 $\lambda = 5e-3$ 
'''
l2_lambda = 5e-3
#####

print ("\nTraining with SGD plus L2 Regularization...")
results_sgd_l2 = train_net(small_data_dict, model_sgd_l2, loss_f_sgd_l2, opt
                           max_epochs=50, show_every=10000, verbose=False, r

opt_params_sgd_l2, loss_hist_sgd_l2, train_acc_hist_sgd_l2, val_acc_hist_sgd

plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgdw, 'o', label="SGD with Weight Decay")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgdw, '-o', label="SGD with Weight Decay")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgdw, '-o', label="SGD with Weight Decay")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd_l1, 'o', label="SGD with L1 Regularization")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd_l1, '-o', label="SGD with L1 Regularization")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd_l1, '-o', label="SGD with L1 Regularization")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd_l2, 'o', label="SGD with L2 Regularization")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd_l2, '-o', label="SGD with L2 Regularization")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd_l2, '-o', label="SGD with L2 Regularization")

```


```
for i in [1, 2, 3]:  
    plt.subplot(3, 1, i)  
    plt.legend(loc='upper center', ncol=4)  
plt.gcf().set_size_inches(15, 15)  
plt.show()
```

Training with SGD plus L2 Regularization...

```
100%|██████████| 200/200 [00:04<00:00, 44.54it/s]
100%|██████████| 200/200 [00:04<00:00, 45.53it/s]
100%|██████████| 200/200 [00:04<00:00, 42.85it/s]
100%|██████████| 200/200 [00:04<00:00, 42.72it/s]
100%|██████████| 200/200 [00:04<00:00, 42.77it/s]
100%|██████████| 200/200 [00:04<00:00, 43.14it/s]
100%|██████████| 200/200 [00:04<00:00, 42.44it/s]
100%|██████████| 200/200 [00:04<00:00, 44.73it/s]
100%|██████████| 200/200 [00:04<00:00, 43.44it/s]
100%|██████████| 200/200 [00:04<00:00, 42.80it/s]
100%|██████████| 200/200 [00:04<00:00, 42.83it/s]
100%|██████████| 200/200 [00:04<00:00, 44.68it/s]
100%|██████████| 200/200 [00:04<00:00, 43.28it/s]
100%|██████████| 200/200 [00:04<00:00, 44.07it/s]
100%|██████████| 200/200 [00:04<00:00, 43.89it/s]
100%|██████████| 200/200 [00:04<00:00, 43.10it/s]
100%|██████████| 200/200 [00:04<00:00, 44.55it/s]
100%|██████████| 200/200 [00:04<00:00, 45.18it/s]
```

[illegible]

```
<00:00, 42.14it/s]
100%|██████████| 200/200 [00:04<
<00:00, 36.06it/s]
100%|██████████| 200/200 [00:05<
<00:00, 39.50it/s]
100%|██████████| 200/200 [00:05<
<00:00, 44.88it/s]
100%|██████████| 200/200 [00:05<
<00:00, 38.94it/s]
100%|██████████| 200/200 [00:05<
<00:00, 37.88it/s]
100%|██████████| 200/200 [00:05<
<00:00, 37.89it/s]
100%|██████████| 200/200 [00:04<
<00:00, 40.35it/s]
100%|██████████| 200/200 [00:05<
<00:00, 36.38it/s]
100%|██████████| 200/200 [00:05<
<00:00, 39.27it/s]
100%|██████████| 200/200 [00:04<
<00:00, 43.90it/s]
100%|██████████| 200/200 [00:05<
<00:00, 37.62it/s]
100%|██████████| 200/200 [00:05<
<00:00, 36.70it/s]
```



Adam [2pt]

The update rule of Adam is as shown below:

$$\begin{aligned}t &= t + 1 \\g_t &: \text{gradients at update step } t \\m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\\hat{m}_t &= m_t / (1 - \beta_1^t) \\\hat{v}_t &= v_t / (1 - \beta_2^t) \\\theta_{t+1} &= \theta_t - \frac{\eta \hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}\end{aligned}$$

Complete the `Adam()` function in `lib/optim.py` Important Notes:

1. t must be updated before everything else
2. β_1^t is β_1 exponentiated to the t 'th power
3. You should also enable weight decay in Adam, similar to what you did in SGD

```

In [382]: %reload_ext autoreload

seed = 1234
np.random.seed(seed=seed)

# Test Adam implementation; you should see errors around 1e-7 or less
N, D = 4, 5
test_adam = sequential(fc(N, D, name="adam_fc"))

w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
m = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)
v = np.linspace(0.7, 0.5, num=N*D).reshape(N, D)

test_adam.layers[0].params = {"adam_fc_w": w}
test_adam.layers[0].grads = {"adam_fc_w": dw}

opt_adam = Adam(test_adam, 1e-2, 0.9, 0.999, t=5)
opt_adam.mt = {"adam_fc_w": m}
opt_adam.vt = {"adam_fc_w": v}
opt_adam.step()

updated_w = test_adam.layers[0].params["adam_fc_w"]
mt = opt_adam.mt["adam_fc_w"]
vt = opt_adam.vt["adam_fc_w"]

expected_updated_w = np.asarray([
    [-0.40094747, -0.34836187, -0.29577703, -0.24319299, -0.19060977],
    [-0.1380274, -0.08544591, -0.03286534, 0.01971428, 0.0722929],
    [ 0.1248705, 0.17744702, 0.23002243, 0.28259667, 0.33516969],
    [ 0.38774145, 0.44031188, 0.49288093, 0.54544852, 0.59801459]])
expected_v = np.asarray([
    [ 0.69966, 0.68908382, 0.67851319, 0.66794809, 0.65738853,],
    [ 0.64683452, 0.63628604, 0.6257431, 0.61520571, 0.60467385,],
    [ 0.59414753, 0.58362676, 0.57311152, 0.56260183, 0.55209767,],
    [ 0.54159906, 0.53110598, 0.52061845, 0.51013645, 0.49966, ]])
expected_m = np.asarray([
    [ 0.48, 0.49947368, 0.51894737, 0.53842105, 0.55789474],
    [ 0.57736842, 0.59684211, 0.61631579, 0.63578947, 0.65526316],
    [ 0.67473684, 0.69421053, 0.71368421, 0.73315789, 0.75263158],
    [ 0.77210526, 0.79157895, 0.81105263, 0.83052632, 0.85 ]])

print ('The following errors should be around or less than 1e-7')
print ('updated_w error: ', rel_error(expected_updated_w, updated_w))
print ('mt error: ', rel_error(expected_m, mt))
print ('vt error: ', rel_error(expected_v, vt))

```

```

The following errors should be around or less than 1e-7
updated_w error: 1.1395691798535431e-07
mt error: 4.214963193114416e-09
vt error: 4.208314038113071e-09

```

Comparing the Weight Decay v.s. L2 Regularization in Adam [5pt]

Run the following code block to compare the plotted results between effects of weight decay and L2 regularization on Adam. Are they still the same? (we can make them the same as in SGD, can we also do it in Adam?)

In [383...

```

seed = 1234
reset_seed(seed)
model_adam_wd = FullyConnectedNetwork()
loss_f_adam_wd = cross_entropy()
optimizer_adam_wd = Adam(model_adam_wd.net, lr=1e-4, weight_decay=1e-6)

print ("Training with AdamW...")
results_adam_wd = train_net(small_data_dict, model_adam_wd, loss_f_adam_wd,
                           max_epochs=50, show_every=10000, verbose=False)

reset_seed(seed)
model_adam_l2 = FullyConnectedNetwork()
loss_f_adam_l2 = cross_entropy()
optimizer_adam_l2 = Adam(model_adam_l2.net, lr=1e-4)
reg_lambda_l2 = 1e-4
print ("\nTraining with Adam + L2...")
results_adam_l2 = train_net(small_data_dict, model_adam_l2, loss_f_adam_l2,
                           max_epochs=50, show_every=10000, verbose=False, reg

opt_params_adam_wd, loss_hist_adam_wd, train_acc_hist_adam_wd, val_acc_hist_
opt_params_adam_l2, loss_hist_adam_l2, train_acc_hist_adam_l2, val_acc_hist_

plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd, 'o', label="Vanilla SGD")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd, '-o', label="Vanilla SGD")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd, '-o', label="Vanilla SGD")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgdw, 'o', label="SGD with Weight Decay")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgdw, '-o', label="SGD with Weight Decay")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgdw, '-o', label="SGD with Weight Decay")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_adam_wd, 'o', label="Adam with Weight Decay")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_adam_wd, '-o', label="Adam with Weight Decay")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_adam_wd, '-o', label="Adam with Weight Decay")

```

```
plt.subplot(3, 1, 1)
plt.plot(loss_hist_adam_l2, 'o', label="Adam with L2")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_adam_l2, '-o', label="Adam with L2")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_adam_l2, '-o', label="Adam with L2")


for i in [1, 2, 3]:
    plt.subplot(3, 1, i)
    plt.legend(loc='upper center', ncol=4)
plt.gcf().set_size_inches(15, 15)
plt.show()
```

Training with AdamW...

```
100%|██████████| 200/200 [00:07<00:00, 26.30it/s]
100%|██████████| 200/200 [00:05<00:00, 37.82it/s]
100%|██████████| 200/200 [00:05<00:00, 36.28it/s]
100%|██████████| 200/200 [00:05<00:00, 35.68it/s]
100%|██████████| 200/200 [00:05<00:00, 33.71it/s]
100%|██████████| 200/200 [00:05<00:00, 37.88it/s]
100%|██████████| 200/200 [00:05<00:00, 36.89it/s]
100%|██████████| 200/200 [00:06<00:00, 28.69it/s]
100%|██████████| 200/200 [00:07<00:00, 27.33it/s]
100%|██████████| 200/200 [00:05<00:00, 37.78it/s]
100%|██████████| 200/200 [00:05<00:00, 34.05it/s]
100%|██████████| 200/200 [00:05<00:00, 35.47it/s]
100%|██████████| 200/200 [00:05<00:00, 37.92it/s]
100%|██████████| 200/200 [00:04<00:00, 44.19it/s]
100%|██████████| 200/200 [00:05<00:00, 39.64it/s]
100%|██████████| 200/200 [00:05<00:00, 38.22it/s]
100%|██████████| 200/200 [00:05<00:00, 38.58it/s]
100%|██████████| 200/200 [00:05<00:00, 39.44it/s]
```


[illegible]

Training with Adam + L2...



[illegible]

[illegible]



Adam + weight decay and Adam + L2 are slightly different since the L2 terms are affected by Adam's momentum whereas regular weight decay is not affected by momentum terms. (L2 updates by $lr/lambdaparams$ whereas weight decay updates by $decay*params$) They can be modified to become them same by adjusting L2's lambda term to account for m_hat and v_hat .

Submission

Please prepare a PDF document `problem_1_solution.pdf` in the root directory of this repository with all plots and inline answers of your solution. Concretely, the document should contain the following items in strict order:

1. Training loss / accuracy curves for the simple neural network training with $> 30\%$ validation accuracy
2. Plots for comparing vanilla SGD to SGD + Weight Decay, SGD + L1 and SGD + L2
3. "Comparing different Regularizations with Adam" plots

Note that you still need to submit the jupyter notebook with all generated solutions. We will randomly pick submissions and check that the plots in the PDF and in the notebook are equivalent.

Problem 2: Incorporating CNNs

- Learning Objective: In this problem, you will learn how to deeply understand how Convolutional Neural Networks work by implementing one.
- Provided Code: We provide the skeletons of classes you need to complete. Forward checking and gradient checkings are provided for verifying your implementation as well.
- TODOs: you will implement a Convolutional Layer and a MaxPooling Layer to improve on your classification results in part 1.

```
In [1]: from lib.mlp.fully_conn import *
        from lib.mlp.layer_utils import *
        from lib.mlp.train import *
        from lib.cnn.layer_utils import *
        from lib.cnn.cnn_models import *
        from lib.datasets import *
        from lib.grad_check import *
        from lib.optim import *
        import numpy as np
        import matplotlib.pyplot as plt

        %matplotlib inline
        plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
        plt.rcParams['image.interpolation'] = 'nearest'
        plt.rcParams['image.cmap'] = 'gray'

        # for auto-reloading external modules
        # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
        %load_ext autoreload
        %autoreload 2
```

Loading the data (CIFAR-100 with 20 superclasses)

In this homework, we will be classifying images from the CIFAR-100 dataset into the 20 superclasses. More information about the CIFAR-100 dataset and the 20 superclasses can be found [here](#).


Download the CIFAR-100 data files [here](#), and save the `.mat` files to the `data/cifar100` directory.

```
In [2]: data = CIFAR100_data('data/cifar100/')
for k, v in data.items():
    if type(v) == np.ndarray:
        print ("Name: {} Shape: {}, {}".format(k, v.shape, type(v)))
    else:
        print("{}: {}".format(k, v))
label_names = data['label_names']
mean_image = data['mean_image'][0]
std_image = data['std_image'][0]
```

```
Name: data_train Shape: (40000, 32, 32, 3), <class 'numpy.ndarray'>
Name: labels_train Shape: (40000,), <class 'numpy.ndarray'>
Name: data_val Shape: (10000, 32, 32, 3), <class 'numpy.ndarray'>
Name: labels_val Shape: (10000,), <class 'numpy.ndarray'>
Name: data_test Shape: (10000, 32, 32, 3), <class 'numpy.ndarray'>
Name: labels_test Shape: (10000,), <class 'numpy.ndarray'>
label_names: ['aquatic_mammals', 'fish', 'flowers', 'food_containers', 'fruit_and_vegetables', 'household_electrical_devices', 'household_furniture', 'insects', 'large_carnivores', 'large_man-made_outdoor_things', 'large_natural_outdoor_scenes', 'large_omnivores_and_herbivores', 'medium_mammals', 'non-insect_invertebrates', 'people', 'reptiles', 'small_mammals', 'trees', 'vehicles_1', 'vehicles_2']
Name: mean_image Shape: (1, 1, 1, 3), <class 'numpy.ndarray'>
Name: std_image Shape: (1, 1, 1, 3), <class 'numpy.ndarray'>
```

```
In [3]: idx = 0
image_data = data['data_train'][idx]
image_data = ((image_data*std_image + mean_image) * 255).astype(np.int32)
plt.imshow(image_data)
label = label_names[data['labels_train'][idx]]
print("Label:", label)
```

```
Label: large_omnivores_and_herbivores
```



Convolutional Neural Networks

We will use convolutional neural networks to try to improve on the results from Problem 1. Convolutional layers make the assumption that local pixels are more important for prediction than far-away pixels. This allows us to form networks that are robust to small changes in positioning in images.

Convolutional Layer Output size calculation [2pts]

As you have learned, two important parameters of a convolutional layer are its stride and padding. To warm up, we will need to calculate the output size of a convolutional layer given its stride and padding. To do this, open the `lib/cnn/layer_utils.py` file and fill out the TODO section in the `get_output_size` function in the `ConvLayer2D` class.

Implement your function so that it returns the correct size as indicated by the block below.

```
In [50]: %reload_ext autoreload

input_image = np.zeros([32, 28, 28, 3]) # a stack of 32 28 by 28 rgb images
in_channels = input_image.shape[-1] #must agree with the last dimension of t
k_size = 4
n_filt = 16

conv_layer = ConvLayer2D(in_channels, k_size, n_filt, stride=2, padding=3)
output_size = conv_layer.get_output_size(input_image.shape)

print("Received {} and expected [32, 16, 16, 16]".format(output_size))

Received (32, 16, 16, 16) and expected [32, 16, 16, 16]
```

Convolutional Layer Forward Pass [5pts]

Now, we will implement the forward pass of a convolutional layer. Fill in the TODO block in the `forward` function of the `ConvLayer2D` class.

```

In [257... %reload_ext autoreload

# Test the convolutional forward function
input_image = np.linspace(-0.1, 0.4, num=1*8*8*1).reshape([1, 8, 8, 1]) # a
in_channels, k_size, n_filt = 1, 5, 2

weight_size = k_size*k_size*in_channels*n_filt
bias_size = n_filt

single_conv = ConvLayer2D(in_channels, k_size, n_filt, stride=1, padding=0,

w = np.linspace(-0.2, 0.2, num=weight_size).reshape(k_size, k_size, in_chann
b = np.linspace(-0.3, 0.3, num=bias_size)

single_conv.params[single_conv.w_name] = w
single_conv.params[single_conv.b_name] = b

out = single_conv.forward(input_image)

print("Received output shape: {}, Expected output shape: (1, 4, 4, 2)".forma

correct_out = np.array([[
    [-0.03874312, 0.57000324],
    [-0.03955296, 0.57081309],
    [-0.04036281, 0.57162293],
    [-0.04117266, 0.57243278]],

    [-0.0452219, 0.57648202],
    [-0.04603175, 0.57729187],
    [-0.04684159, 0.57810172],
    [-0.04765144, 0.57891156]],

    [-0.05170068, 0.5829608 ],
    [-0.05251053, 0.58377065],
    [-0.05332038, 0.5845805 ],
    [-0.05413022, 0.58539035]],

    [-0.05817946, 0.58943959],
    [-0.05898931, 0.59024943],
    [-0.05979916, 0.59105928],
    [-0.06060901, 0.59186913]]]])

# Compare your output with the above pre-computed ones.
# The difference should not be larger than 1e-7
print ("Difference: ", rel_error(out, correct_out))

Received output shape: (1, 4, 4, 2), Expected output shape: (1, 4, 4, 2)
Difference: 5.110565335399418e-08

```

Conv Layer Backward [5pts]

Now complete the backward pass of a convolutional layer. Fill in the TODO block in the `backward` function of the `ConvLayer2D` class. Check you results with this code and expect differences of less than $1e-6$.

```
In [273] %reload_ext autoreload

# Test the conv backward function
img = np.random.randn(15, 8, 8, 3)
w = np.random.randn(4, 4, 3, 12)
b = np.random.randn(12)
dout = np.random.randn(15, 4, 4, 12)

single_conv = ConvLayer2D(input_channels=3, kernel_size=4, number_filters=12)
single_conv.params[single_conv.w_name] = w
single_conv.params[single_conv.b_name] = b

dimg_num = eval_numerical_gradient_array(lambda x: single_conv.forward(img),
dw_num = eval_numerical_gradient_array(lambda w: single_conv.forward(img), w
db_num = eval_numerical_gradient_array(lambda b: single_conv.forward(img), b

out = single_conv.forward(img)

dimg = single_conv.backward(dout)
dw = single_conv.grads[single_conv.w_name]
db = single_conv.grads[single_conv.b_name]

# The error should be around 1e-6
print("dimg Error: ", rel_error(dimg_num, dimg))
# The errors should be around 1e-8
print("dw Error: ", rel_error(dw_num, dw))
print("db Error: ", rel_error(db_num, db))
# The shapes should be same
print("dimg Shape: ", dimg.shape, img.shape)

dimg Error:  4.644590234305012e-08
dw Error:   3.127392012612529e-09
db Error:   3.7320812919197414e-10
dimg Shape:  (15, 8, 8, 3) (15, 8, 8, 3)
```

Max pooling Layer

Now we will implement maxpooling layers, which can help to reduce the image size while preserving the overall structure of the image.

Forward Pass max pooling [5pts]

Fill out the TODO block in the `forward` function of the `MaxPoolingLayer` class.

```
In [274... # Test the convolutional forward function
input_image = np.linspace(-0.1, 0.4, num=64).reshape([1, 8, 8, 1]) # a single channel image

maxpool= MaxPoolingLayer(pool_size=4, stride=2, name="maxpool_test")
out = maxpool.forward(input_image)

print("Received output shape: {}, Expected output shape: (1, 3, 3, 1)".format(out.shape))

correct_out = np.array([[
    [0.11428571],
    [0.13015873],
    [0.14603175]],

    [[0.24126984],
    [0.25714286],
    [0.27301587]],

    [[0.36825397],
    [0.38412698],
    [0.4         ]]])

# Compare your output with the above pre-computed ones.
# The difference should not be larger than 1e-7
print ("Difference: ", rel_error(out, correct_out))

Received output shape: (1, 3, 3, 1), Expected output shape: (1, 3, 3, 1)
Difference: 1.8750000280978013e-08
```

Backward Pass Max pooling [5pts]

Fill out the `backward` function in the `MaxPoolingLayer` class.

```
In [275... img = np.random.randn(15, 8, 8, 3)

dout = np.random.randn(15, 3, 3, 3)

maxpool= MaxPoolingLayer(pool_size=4, stride=2, name="maxpool_test")

dimg_num = eval_numerical_gradient_array(lambda x: maxpool.forward(img), img, dout)

out = maxpool.forward(img)
dimg = maxpool.backward(dout)

# The error should be around 1e-8
print("dimg Error: ", rel_error(dimg_num, dimg))
# The shapes should be same
print("dimg Shape: ", dimg.shape, img.shape)

dimg Error: 3.279406847766653e-12
dimg Shape: (15, 8, 8, 3) (15, 8, 8, 3)
```

Test a Small Convolutional Neural Network [3pts]

Please find the `TestCNN` class in `lib/cnn/cnn_models.py`. Again you only need to complete few lines of code in the TODO block. Please design a Convolutional --> Maxpool --> flatten --> fc network where the shapes of parameters match the given shapes. Please insert the corresponding names you defined for each layer to `param_name_w`, and `param_name_b` respectively. Here you only modify the `param_name` part, the `_w`, and `_b` are automatically assigned during network setup.

```

In [287... %reload_ext autoreload

seed = 1234
np.random.seed(seed=seed)

model = TestCNN()
loss_func = cross_entropy()

B, H, W, iC = 4, 8, 8, 3 #batch, height, width, in_channels
k = 3 #kernel size
oC, Hi, O = 3, 27, 5 # out channels, Hidden Layer input, Output size
std = 0.02
x = np.random.randn(B,H,W,iC)
y = np.random.randint(0, size=B)

print ("Testing initialization ... ")

#####
# TODO: param_name should be replaced accordingly #
#####
w1_std = abs(model.net.get_params("conv1_w").std() - std)
b1 = model.net.get_params("conv1_b").std()
w2_std = abs(model.net.get_params("fc1_w").std() - std)
b2 = model.net.get_params("fc1_b").std()
#####
#                               END OF YOUR CODE                               #
#####

assert w1_std < std / 10, "First layer weights do not seem right"
assert np.all(b1 == 0), "First layer biases do not seem right"
assert w2_std < std / 10, "Second layer weights do not seem right"
assert np.all(b2 == 0), "Second layer biases do not seem right"
print ("Passed!")

print ("Testing test-time forward pass ... ")
w1 = np.linspace(-0.7, 0.3, num=k*k*iC*oC).reshape(k,k,iC,oC)
w2 = np.linspace(-0.2, 0.2, num=Hi*O).reshape(Hi, O)
b1 = np.linspace(-0.6, 0.2, num=oC)
b2 = np.linspace(-0.9, 0.1, num=O)

#####
# TODO: param_name should be replaced accordingly #
#####
model.net.assign("conv1_w", w1)
model.net.assign("conv1_b", b1)
model.net.assign("fc1_w", w2)
model.net.assign("fc1_b", b2)
#####
#                               END OF YOUR CODE                               #
#####

feats = np.linspace(-5.5, 4.5, num=B*H*W*iC).reshape(B,H,W,iC)
scores = model.forward(feats)
print(scores)
correct_scores = np.asarray([[-13.85107294, -11.52845818, -9.20584342, -6.
[-11.44514171, -10.21200524, -8.97886878, -7.74573231, -6.51259584],

```

```

[ -9.03921048, -8.89555231, -8.75189413, -8.60823596, -8.46457778],
[ -6.63327925, -7.57909937, -8.52491949, -9.4707396, -10.41655972]])
scores_diff = np.sum(np.abs(scores - correct_scores))
assert scores_diff < 1e-6, "Your implementation might be wrong!"
print ("Passed!")

print ("Testing the loss ...",)
y = np.asarray([0, 2, 1, 4])
loss = loss_func.forward(scores, y)
dLoss = loss_func.backward()
correct_loss = 4.56046848799693
assert abs(loss - correct_loss) < 1e-10, "Your implementation might be wrong"
print ("Passed!")

print ("Testing the gradients (error should be no larger than 1e-6) ...")
din = model.backward(dLoss)
for layer in model.net.layers:
    if not layer.params:
        continue
    for name in sorted(layer.grads):
        f = lambda _: loss_func.forward(model.forward(feats), y)
        grad_num = eval_numerical_gradient(f, layer.params[name], verbose=False)
        print ('%s relative error: %.2e' % (name, rel_error(grad_num, layer.

Testing initialization ...
Passed!
Testing test-time forward pass ...
[[-13.85107294 -11.52845818 -9.20584342 -6.88322866 -4.5606139 ]
 [-11.44514171 -10.21200524 -8.97886878 -7.74573231 -6.51259584]
 [ -9.03921048 -8.89555231 -8.75189413 -8.60823596 -8.46457778]
 [ -6.63327925 -7.57909937 -8.52491949 -9.4707396 -10.41655972]]
5.304109507164867e-08
Passed!
Testing the loss ...
Passed!
Testing the gradients (error should be no larger than 1e-6) ...
conv1_b relative error: 4.88e-09
conv1_w relative error: 9.26e-10
fc1_b relative error: 8.77e-11
fc1_w relative error: 3.89e-07

```

Training the Network [25pts]

In this section, we defined a `SmallConvolutionalNetwork` class for you to fill in the TODO block in `lib/cnn/cnn_models.py`.

Here please design a network with at most two convolutions and two maxpooling layers (you may use less). You can adjust the parameters for any layer, and include layers other than those listed above that you have implemented (such as fully-connected layers and non-linearities). You are also free to select any optimizer you have implemented (with any learning rate).

You will train your network on CIFAR-100 20-way superclass classification. Try to find a combination that is able to achieve 40% validation accuracy.

Since the CNN takes significantly longer to train than the fully connected network, it is suggested to start off with fewer filters in your Conv layers and fewer intermediate fully-connected layers so as to get faster initial results.

In [288... *# Arrange the data*

```
data_dict = {
    "data_train": (data["data_train"], data["labels_train"]),
    "data_val": (data["data_val"], data["labels_val"]),
    "data_test": (data["data_test"], data["labels_test"])
}
```

In [289...

```
print("Data shape:", data_dict["data_train"][0].shape)
print("Flattened data input size:", np.prod(data["data_train"].shape[1:]))
print("Number of data classes:", max(data['labels_train']) + 1)
```

```
Data shape: (40000, 32, 32, 3)
Flattened data input size: 3072
Number of data classes: 20
```



```
%reload_ext autoreload

seed = 123
np.random.seed(seed=seed)

model = SmallConvolutionalNetwork()
loss_f = cross_entropy()


results = None
#####
# TODO: Use the train_net function you completed to train a network
# You may only adjust the hyperparameters within this block
#####
optimizer = Adam(model.net, 1e-3)

batch_size = 100#50#10
epochs = 10#5
lr_decay = 1.0#.999
lr_decay_every = 30#10
regularization = "none"
reg_lambda = 0.01
#####
#                               END OF YOUR CODE
#####
results = train_net(data_dict, model, loss_f, optimizer, batch_size, epochs,
                    lr_decay, lr_decay_every, show_every=4000, verbose=True,
                    opt_params, loss_hist, train_acc_hist, val_acc_hist = results
```

[illegible]

Run the code below to generate the training plots.

In [372...



Visualizing Layers [5pts]

An interesting finding from early research in convolutional networks was that the learned convolutions resembled filters used for things like edge detection. Complete the code below to visualize the filters in the first convolutional layer of your best model.

```
In [373]: im_array = None
nrows, ncols = None, None

#####
# TODO: read the weights in the convolutional      #
# layer and reshape them to a grid of images to   #
# view with matplotlib.                           #
#####
layer1_params = model.net.get_params('conv1_w')
im_array = np.zeros(shape=(layer1_params.shape[:-1]))


for i in range(0, layer1_params.shape[3]):
    if i == 0:
        im_array = layer1_params[:, :, :, i]
    else:
        im_array = np.append(im_array, layer1_params[:, :, :, i], axis=0)

# https://stackoverflow.com/questions/49922460/scale-a-numpy-array-with-from
im_array = ((im_array - im_array.min()) * (1/(im_array.max() - im_array.min(

#####
#                               END OF YOUR CODE                               #
#####

plt.imshow(im_array)
```

```
Out[373]: <matplotlib.image.AxesImage at 0x7ff68ab091f0>
```



Inline Question: Comment below on what kinds of filters you see. Include your response in your submission [5pts]

The RGB weights learned usually correspond to important local components in the original image. It isn't

clear from above since the a larger kernel size could have been used for a 5x5 kernel convolved over a 32x32 image

Extra-Credit: Analysis on Trained Model [5pts]

For extra credit, you can perform some additional analysis of your trained model. Some suggested analyses are:

1. Plot the [confusion matrix](#) of your model's predictions on the test set. Look for trends to see which classes are frequently misclassified as other classes (e.g. are the two vehicle superclasses frequently confused with each other?).
2. Implement [BatchNorm](#) and analyze how the models train with and without BatchNorm.
3. Introduce some small noise in the labels, and investigate how that affects training and validation accuracy.

You are free to choose any analysis question of interest to you. We will not be providing any starter code for the extra credit. Include your extra-credit analysis as the final section of your report pdf, titled "Extra Credit".

Submission

Please prepare a PDF document `problem_2_solution.pdf` in the root directory of this repository with all plots and inline answers of your solution. Concretely, the document should contain the following items in strict order:

1. Training loss / accuracy curves for CNN training
2. Visualization of convolutional filters
3. Answers to inline questions about convolutional filters

Note that you still need to submit the jupyter notebook with all generated solutions. We will randomly pick submissions and check that the plots in the PDF and in the notebook are equivalent.

In []: